

2015 Computational Social Science Workshop

Day 1 - Introduction to `python` - Part 1 / 3

with **Jongbin Jung** (jongbin at stanford.edu)

- PhD student. Decision analysis, MS&E

All material for days 1 (intro to `python`) and 2 (web scraping with `python`) publicly available [here \(https://github.com/jongbinjung/css-python-workshop\)](https://github.com/jongbinjung/css-python-workshop)

Meet `python`

Why `python`?

- fast
- modular
- object oriented
- tons of libraries
- (relatively) easy to read
- *gets the job done*

Choice of primary programming language should really depend on your preferred style of thinking and tools of choice (e.g., linear models? decision trees?). But you'll never know if it's right (or wrong) for you unless you give it a try!

What do you mean by python?

Python is the name of a programming language. Period. But different people could have different ideas about what 'python' looks like. Three broad categories I have in mind are:

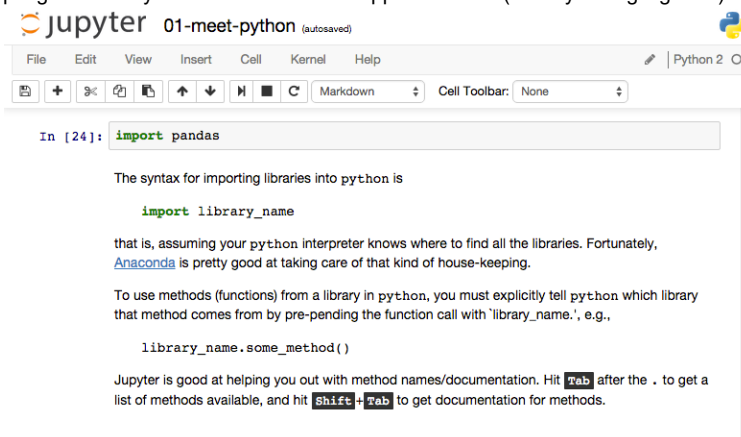
1. Running interactive commands in the python interpreter (aka, the glorified calculator)

```
Python 2.7.10 [Anaconda 2.3.0 (x86_64)] (default, May 28 2015, 17:04:42)
[GCC 4.2.1 (Apple Inc. build 5577)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
Anaconda is brought to you by Continuum Analytics.
Please check out: http://continuum.io/thanks and https://binstar.org
>>> x, y = 5, 9
>>> x + y
14
>>> |
```

2. Python development in some kind of text editor or specialized environment

```
1 import argparse
2 import pandas
3 import os
4 import os.path
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.tree import DecisionTreeClassifier
7
8 # set up argument parser
9 parser = argparse.ArgumentParser(
10     description='Predict using single decision tree and random forest on ' +
11     'cleaned/parsed dataset')
12 parser.add_argument('-d', dest='data_dir', required=True,
13     help='path to data directory')
14 parser.add_argument('-njobs', default=20, type=int,
15     help='number of parallel jobs to run (default 20)')
16 parser.add_argument('-s', dest='tree_size', default=0, type=int,
17     help='number of trees in the forest (default njobs*100)')
18 parser.add_argument('-data', dest='data_file', default='parsed folds.tsv.gz',
19     type=str, help='file name for parsed data')
20
21 args = parser.parse_args()
22 # set the relative data directory
23
24 data_mnt = args.data_dir
25
26 if data_mnt[-1] != '/':
27     data_mnt += '/'
28
```

3. Research-type scripting with heavy documentation and snippets of code (usually wrangling data)



The screenshot shows a Jupyter Notebook window titled "01-meet-python (autosaved)". The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Help) and a toolbar with icons for running, saving, and other actions. The current cell is a code cell containing the command `import pandas`. Below the code cell, the notebook displays a text explanation of the `import` statement, including examples and instructions on how to use methods from a library.

```
In [24]: import pandas
```

The syntax for importing libraries into python is

```
import library_name
```

that is, assuming your python interpreter knows where to find all the libraries. Fortunately, [Anaconda](#) is pretty good at taking care of that kind of house-keeping.

To use methods (functions) from a library in python, you must explicitly tell python which library that method comes from by pre-pending the function call with 'library_name.', e.g.,

```
library_name.some_method()
```

Jupyter is good at helping you out with method names/documentation. Hit `Tab` after the `.` to get a list of methods available, and hit `Shift+Tab` to get documentation for methods.

For the purpose of a proper "python introduction", we'll look at all three flavors, and try to spend more time on the workflow that participants find most relevant to their work/research.

(Some of the material for Parts 1 and 2 were adapted from [The Python Tutorial](https://docs.python.org/2/tutorial/) (<https://docs.python.org/2/tutorial/>))

Anaconda

- **Anaconda** (<https://store.continuum.io/cshop/anaconda/>) is one of many python distributions
- it's free (some distributions aren't).
- comes pre-installed with a lot of useful libraries
- support many different workflows out-of-the-box
- think of it as a "**python starter kit**"
- as far as I've seen, it's the most painless way to get things up-and-running
- download/install from [here](https://store.continuum.io/cshop/anaconda/) (<https://store.continuum.io/cshop/anaconda/>)

1. python interpreter (aka glorified calculator)

You can start using python right away, by running the python interpreter in the command line (Windows) or terminal (OS X / *nix)

- **Windows**
 - Win + r, type python and hit OK (or Enter)
- **OS X / *nix**
 - Open a terminal
 - type python and hit Enter

```
$ python
Python 2.7.10 [Anaconda 2.3.0 (x86_64)] (default, May 28 2015, 17:04:42)
[GCC 4.2.1 (Apple Inc. build 5577)] on darwin
Type "help", "copyright", "credits" or "license" for more information.
Anaconda is brought to you by Continuum Analytics.
Please check out: http://continuum.io/thanks and https://binstar.org
>>> |
```

Numbers

Most of your numerical needs will fall into two types: integers (1, 2, ...) and floating point (0.5, 1.75, ...). Let's start with some basic calculations:

```
In [1]: 296719 + 445567 # add stuff
```

```
Out[1]: 742286
```

```
In [2]: 762301 - 162379 # subtract stuff
```

```
Out[2]: 599922
```

```
In [3]: 148501 * 220131 # multiply stuff
```

```
Out[3]: 32689673631L
```

```
In [4]: 1 / 10 # divide stuff ...?
```

```
Out[4]: 0
```

You might notice two things:

1. Anything after the # is ignored. These are called **comments**. It's good practice to use them to leave notes for *future* you.
2. 1/10 is **not** 0.

What's happening?

Python never asks you what type each number/variable is. It assumes the type of a variable based on its value, in this case 1 and 10 are both assumed to be integers. And the division operation (/) determines the type of its result based on the type of the operands. If both operands are integers, *floor division* is performed and an integer is returned. If either operand is a float, classic division is performed and a float is returned.

So, to *implicitly* tell python that you're expecting a floating point number, you have to make one (or both) of the numbers look like a floating point number:

```
In [5]: 1 / 10.0
```

```
Out[5]: 0.1
```

This is an important thing to know about python, especially if you're used to different programming languages (statically typed languages such as C, Java, ...)

Python will *not* keep track of types for you.

It is absolutely up to you to make sure that your variables are the right type.

Some more useful operators are:

```
In [6]: 5.0 // 2 # an explicit floor division
```

```
Out[6]: 2.0
```

```
In [7]: 5.0 % 2 # modulo operation (finds the remainder after division)
```

```
Out[7]: 1.0
```

```
In [8]: 2 ** 10 # calculate powers
```

```
Out[8]: 1024
```

You can assign (store) values to variables with an equal (=) sign. Again, since python isn't interested in knowing the type of each variable before runtime, you don't need to declare any. Many people love or hate python for this.

```
In [9]: sales = 1000
        costs = 1300
        sales - costs
```

```
Out[9]: -300
```

Note that the interpreter doesn't return anything when you assign/define variables. Also, trying to use variables without defining them will result in an error.

```
In [10]: profit
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-10-65280be1bd6e> in <module>()  
----> 1 profit  
  
NameError: name 'profit' is not defined
```

There are more types (e.g., long - for huge integers, complex numbers, fractions) if you think you'll need them. See [the docs](https://docs.python.org/2/library/types.html) (<https://docs.python.org/2/library/types.html>) for a comprehensive documentation of python's built-in types.

Strings

One of the many things that python is *really* good for is manipulating strings. Strings are enclosed in either single quotes('...') or double quotes("..."). The only difference between the two is that you need to escape literal double quotes with \ if you're using single quotes, and vice versa. Examples,

```
In [11]: 'This is a string in single quotes, and it works fine'
```

```
Out[11]: 'This is a string in single quotes, and it works fine'
```

```
In [12]: 'This is a string in single quotes, but it's broken!'
```

```
File "<ipython-input-12-191c80bf77a1>", line 1  
    'This is a string in single quotes, but it's broken!'  
                                         ^  
SyntaxError: invalid syntax
```

```
In [13]: 'This is a string with single quotes, and now it\'s fixed!'
```

```
Out[13]: "This is a string with single quotes, and now it's fixed!"
```

```
In [14]: "So, there's incentive to use double quotes"
```

```
Out[14]: "So, there's incentive to use double quotes"
```

That should make it pretty clear. Note that when escaping a single quote in a string enclosed in single quotes, the interpreter internally changes the enclosing quotes to double quotes. The two are absolutely identical, so you should use whichever set of quotes you prefer to enclose your strings, as long as you're consistent.

Use the `print` statement to make output more readable by omitting the enclosing quotes and printing special characters escaped with \.

```
In [15]: 'enclosing double quotes (") and single quotes(\'\') are the same thing in python'
```

```
Out[15]: 'enclosing double quotes (") and single quotes(\'\') are the same thing in python  
,
```

```
In [16]: print 'enclosing double quotes (") and single quotes(\'\') are the same thing in py  
thon'
```

```
enclosing double quotes (") and single quotes(') are the same thing in python
```

Strings can also have more than one line. You can either use an explicit line-break character (\n):

```
In [17]: print 'This string has\ntwo lines!'
```

```
This string has  
two lines!
```

... or use triple quotes `'''...'''` or `"""..."""`

```
In [18]: print '''  
This string has  
two lines!  
'''
```

```
This string has  
two lines!
```

If you look carefully enough, you'll notice that the last string actually has four lines. This is because the triple quotes literally encode all white spaces, including the new lines after the first `'''` and the last `'''`. To avoid this, you can escape the new lines with `\`

```
In [19]: print '''\  
This string (really) has  
two lines!\  
'''
```

```
This string (really) has  
two lines!
```

Unlike some other stingy languages, the plus operator (+) does exactly what you'd expect it to do with strings!

```
In [20]: first_name = 'Jongbin'  
last_name = 'Jung'  
full_name = first_name + ' ' + last_name  
print 'Hello,', full_name
```

```
Hello, Jongbin Jung
```

Even the multiplying operator (*) works!

```
In [21]: print 'Sing, ' + 'la ' * 3
```

```
Sing, la la la
```

Strings, like many things in python, can be indexed (subscripted). The first element (character) has index 0.

```
In [22]: job = 'jedi'  
job[0] # character at position 0
```

```
Out[22]: 'j'
```

Python will yell at you if you go out of range.

```
In [23]: job[10]
```

```
-----  
IndexError                                Traceback (most recent call last)  
<ipython-input-23-72005d25ec5c> in <module>()  
----> 1 job[10]  
  
IndexError: string index out of range
```

But you *can* go backwards with a negative index, -1 being the right-most character.

```
In [24]: job[-1]  # right-most character
```

```
Out[24]: 'i'
```

Slicing is another useful way to get subsets of your string.

```
In [25]: job[0:2]  # characters from position 0 (included) to 2 (excluded)
```

```
Out[25]: 'je'
```

```
In [26]: job[2:4]  # characters from position 2 (included) to 3 (excluded)
```

```
Out[26]: 'di'
```

Omitting the first slice index will default to zero (the first element)

```
In [27]: job[:3]  # slice from the first element (included) to position 3
```

```
Out[27]: 'jed'
```

Omitting the second index will default to the size of the string, which includes the last element.

```
In [28]: job[1:]  # slice from position 1 (included) to the end
```

```
Out[28]: 'edi'
```

Use slices creatively, to make your life easier!

```
In [29]: job[-2:]  # slice the last two characters
```

```
Out[29]: 'di'
```

Unlike indexing, slicing is generous to ambitious ranges

```
In [30]: job[0:100]
```

```
Out[30]: 'jedi'
```

Strings in python are `immutable`, meaning they can't be changed. In other words, you can't assign a value to a string index.

```
In [31]: job[0] = 'T'
```

```
-----  
TypeError                                Traceback (most recent call last)  
<ipython-input-31-d4729ccf68c2> in <module>()  
----> 1 job[0] = 'T'  
  
TypeError: 'str' object does not support item assignment
```

Instead, you have to build a new string from the existing string.

```
In [32]: job = 'T' + job[1:]  
print job  
  
Tedi
```

The built-in `len()` function will return the length of a string.

```
In [33]: len(job)  
  
Out[33]: 4
```

Lists

Lists are perhaps the most versatile array-like data types in python. A list is declared (and printed) as a comma-separated value between square brackets.

```
In [34]: names = ['Luke', 'Leia', 'Han']  
print names  
  
['Luke', 'Leia', 'Han']
```

```
In [35]: numbers = [45, 8, 21, 30]  
print numbers  
  
[45, 8, 21, 30]
```

```
In [36]: mixed = [3, 4, 'five', 'six']  
print mixed  
  
[3, 4, 'five', 'six']
```

Lists (and in fact, all built-in python [sequence](https://docs.python.org/2/glossary.html#term-sequence) (<https://docs.python.org/2/glossary.html#term-sequence>) types) can be indexed and sliced.

```
In [37]: names[0] # returns the item at position 0  
  
Out[37]: 'Luke'
```

```
In [38]: mixed[1:3] # slicing returns a new list  
  
Out[38]: [4, 'five']
```

Also like strings, the `+` operation concatenates two lists


```
In [39]: names + ['Anakin', 'Padme']  
Out[39]: ['Luke', 'Leia', 'Han', 'Anakin', 'Padme']
```

Unlike strings, lists are mutable, which means you can change item values, even to different types.

```
In [40]: mixed[0] = 30 # change indexed item  
print mixed  
[30, 4, 'five', 'six']
```

You can use the `append()` method to add new items to a list. Methods of an object are called directly on the object with a dot (`.`)

```
In [41]: mixed.append('seven')  
print mixed  
[30, 4, 'five', 'six', 'seven']
```

You can also assign values to slices of a list.

```
In [42]: mixed[0:2] = ['three', 'four'] # change sliced list items  
print mixed  
['three', 'four', 'five', 'six', 'seven']
```

Because assignment to a slice replaces a subset of the list, it can be used to delete items.

```
In [43]: mixed[2:] = [] # delete items from position 2 to the end of the list  
print mixed  
['three', 'four']
```

The built-in `len()` function also works as it would with strings.

```
In [44]: len(names)  
Out[44]: 3  
  
In [45]: len(numbers)  
Out[45]: 4  
  
In [46]: len(mixed)  
Out[46]: 2
```

More on Strings and Lists

Both strings and lists have a bunch of helpful methods.

See the docs [here \(https://docs.python.org/2/library/stdtypes.html#string-methods\)](https://docs.python.org/2/library/stdtypes.html#string-methods) for strings and [here \(https://docs.python.org/2/tutorial/datastructures.html#more-on-lists\)](https://docs.python.org/2/tutorial/datastructures.html#more-on-lists) for lists.

Some of the more often used methods are illustrated below:

```

In [47]: s = 'Hi, my name is Jongbin. Yes, Jongbin.'
         s.split() # splits the string into a list (split at spaces by default)

Out[47]: ['Hi,', 'my', 'name', 'is', 'Jongbin.', 'Yes,', 'Jongbin.']

In [48]: s.split(',') # can specify which character to split at

Out[48]: ['Hi', ' my name is Jongbin. Yes', ' Jongbin.']

In [49]: s.count('n') # count the number of non-overlapping occurrences of a substring

Out[49]: 5

In [50]: s.count('Jongbin')

Out[50]: 2

In [51]: s.upper() # makes everything uppercase

Out[51]: 'HI, MY NAME IS JONGBIN. YES, JONGBIN.'

In [52]: s.lower() # makes everything lowercase

Out[52]: 'hi, my name is jongbin. yes, jongbin.'

In [53]: s.lower().count('y') # methods that return a string can be chained

Out[53]: 2

In [54]: ':'.join(s.split())

Out[54]: 'Hi,:my:name:is:Jongbin.:Yes,:Jongbin.'

```

That last one is a little tricky. So, `str.join(some_sequence)` will take each item of `some_sequence` and stick them together with the value of `str` inbetween, making a single large string. It may seem like a crazy thing to do, but is actually pretty useful when converting data into comma-separated values. i.e.,

```

In [55]: ','.join(['some','data','in','a','list'])

Out[55]: 'some,data,in,a,list'

```

And also some tricks with lists:

```

In [56]: fruits = ['mango', 'apple', 'banana', 'apple', 'kiwi', 'mango']
         fruits.append('orange') # add an item to the end of the list
         print fruits

         ['mango', 'apple', 'banana', 'apple', 'kiwi', 'mango', 'orange']

In [57]: fruits.remove('apple') # remove the first item in the list that matches the argument
         print fruits

         ['mango', 'banana', 'apple', 'kiwi', 'mango', 'orange']

In [58]: fruits.index('mango') # return the index of the first item matching the argument

Out[58]: 0

In [59]: fruits.count('mango') # return the number of times x appears in the list

Out[59]: 2

```

```
In [60]: eat = fruits.pop(0) # return and remove item at position 0 from list (removes last
        # item if no index is specified)
        print fruits

        ['banana', 'apple', 'kiwi', 'mango', 'orange']

In [61]: print eat # the item previously at position 0 ('mango') is now 'popped' into the va
        # riable eat

        mango
```

Not exactly a method, but the `in` keyword is useful for checking if a list contains a particular item.

```
In [62]: 'banana' in fruits

Out[62]: True

In [63]: 'tomato' in fruits

Out[63]: False
```

You can still do much more with `python`, even in the interpreter. But anything beyond a single line starts getting tedious in the interpreter, so now might be a good time to move on to the next stage: **python development**.

Exercise 1.

Before moving on, just to make sure that everyone's on the same page, let's try a few tasks:

1. declare a string variable named `s`, that has the value

```
"double quotes" and single 'quotes' are equally acceptable in python
```

2. make `python` count how many times the letter `t` appears in the string `s`
3. replace all quotation marks in the string `s` with an underbar ('_')
 - to be precise, you're not *replacing* the quotations, but *reassigning* the variable `s` with a copy of the old `s` that has underbars replacing quotations; remember that `python` strings are **immutable** (i.e., they are NEVER changed, only reassigned)
 - this was not covered above, but you should use `python`'s `replace()` method for strings; now would be a good time to practice reading the docs (<https://docs.python.org/2/library/stdtypes.html#str.replace> (<https://docs.python.org/2/library/stdtypes.html#str.replace>))
4. split the string `s` into a list named `words`
5. count the length of the string `s` and the list `words`
6. join the elements of the list `words` into a comma separated string